

A Brief History of AI

From Foundations to Generative AI

- Presenter:
- Core idea: to understand today's AI, we must know its **history and shifts**

"A generation which ignores history has no past – and no future."

— Robert A. Heinlein

The Elusive Definition of Intelligence

- “As soon as it works, no one calls it AI anymore.” — John McCarthy
- Intelligence is a **moving target** as technologies become routine.
- Today: AI literacy is no longer only for computer scientists.
- Students in **all disciplines** must:
 - Evaluate when AI is appropriate
 - Defend or challenge model choices for their domain

Roadmap of This Talk

- **Foundations of Computing**
 - Traditional programming vs. data-centric learning
- **The Scaling Problem**
 - Data, storage, and compute infrastructure
- **The AI Landscape**
 - AI vs. Machine Learning vs. Deep Learning
- **Historical Milestones**
 - From early neural nets to Transformers
- **Trust, Safety & Boundaries**
 - Adversarial risk, privacy, and ethics (high-level preview)

SECTION I

Historical Evolution of AI

AI Roots (1940s–1950s)

Foundational ideas:

- **1943 – McCulloch & Pitts**
 - First mathematical model of a **neuron**
 - Basis for artificial neural networks
- **1950 – Alan Turing**
 - Proposes the **Turing Test**:
- **1956 – Dartmouth Workshop**
 - McCarthy, Minsky, Rochester, Shannon
 - Coins the term **"Artificial Intelligence"**




Mapping the AI Landscape

Nested vocabulary (conceptual diagram):

- **Artificial Intelligence**
 - Rule-based systems, search, logic
 -  **Machine Learning**
 - Linear/logistic regression, random forests, SVMs
 -  **Deep Learning**
 - MLPs, CNNs, RNNs, LSTMs, autoencoders
 -  **Generative AI**
 - GANs, Transformers, Large Language Models (LLMs)


| Not all AI is ML; not all ML is deep learning; not all deep learning is generative.

Nested vocabulary (conceptual diagram)



The McCulloch-Pitts Neuron (1943)


- Inspired by neuroscience:
The brain is composed of neurons interacting via synapses, transmitting either excitatory or inhibitory signals
- The neuron can be thought of as a binary classifier using binary input
- A neuron
 - fires if: sum of inputs reaches threshold
 - unless inhibited by other inputs



The Perceptron (1957)

- Frank Rosenblatt's original work on the perceptron in the late 1950s laid the groundwork for subsequent research in neural network theory and applications.

The perceptron model demonstrated the potential of neural networks for pattern recognition tasks



General Purpose Solver (GPS) (1957)

- GPS by Herbert Simon, J.C. Shaw, and Allen Newell
 - a computer program created intended to work as a universal problem solver machine
 - employs heuristic search to explore possible solutions


| ability to reason and make decisions based on logical inference and problem-solving strategies

Expert Systems: Eliza (1966) - short video


 [The First Ever AI Chatbot: ELIZA -1966](#)

 [ELIZA | Artificial Intelligence The Old-School Way | HowTo Install](#)

Expert Systems: Eliza (1966) - short video




Expert Systems: SHRDLU (1968) - short video



The Perceptron (controversy) - Minsky and Papert (1969)

- A single perceptron is unable to compute the XOR function, as it's not linearly separable

Minsky and Papert. "Perceptron: An Introduction to Computational Geometry", book 1969



Expert Systems: Symbolic AI at Scale

- **Dendral (1960s)** : Helped chemists identify molecules from mass spectra.
- **MYCIN (1970s)** : Diagnosed bacterial infections, recommended antibiotics.
- **XCON (late 1970s–1980s)** : Configured VAX computers at DEC; saved tens of millions of dollars.
- **Deep Blue (1990s)** : Defeated world chess champion Garry Kasparov using search + rule-based evaluation.

These systems showed symbolic AI can work very well in narrow, structured domains.

The Multi-layer Perceptron

Training before MLP was challenging and less systematic

- Perceptron Learning Rule
- Heuristic Methods
- Optimization Techniques
- Evolutionary Algorithms


Paul Werbos in 1974 laid the foundation

for training neural networks with **multiple layers of neurons**, which became known as the backpropagation algorithm


Birth of Deep Learning

Deep learning models:

- Stack many **hidden layers** to learn complex features from raw data.
- **Forward pass**
 - Data flows through layers to produce an output
 - Compute loss vs. ground truth
- **Backward pass (backpropagation)**
 - **chain rule** to compute gradients layer by layer
 - Update weights to reduce loss



Evolution of Deep learning



From Hand-Crafted to Learned Features

Evolution in computer vision:


- Early era:
 - Hand-crafted features + probabilistic models (e.g. HMMs for sequences)
- Deep learning era:
 - **Convolutional Neural Networks (CNNs):**
 - Automatically learn spatial patterns (edges, textures, objects)
 - Dramatically outperform manual feature pipelines

CNNs shifted vision from **feature engineering** to **feature learning**.

Recurrent Neural Network (1997)

- A recurrent neural network used in speech recognition and NLP.
- RNNs recognize data's sequential characteristics and use patterns to predict the next likely scenario



Solves the problem - vanishing and exploding gradients



Convolution Neural Network (1998)


LeNet-5 (Yann LeCun)

- Introduced:
 - Local receptive fields
 - Shared weights (convolutions)
 - Subsampling / pooling
- Achieved $\approx 99\%$ accuracy on digit recognition




CNN, RNN, LSTM


(a) Typical ANN



(b) RNN





(c) LSTM



CMNIST dataset (1998)

- Handwritten digits 0–9
- Became a standard benchmark for image recognition




Beyond Digits: CIFAR Datasets

CIFAR-10 / CIFAR-100 (2009)

- Small natural images across many classes:
 - Animals, vehicles, everyday objects
- Posed a harder challenge than MNIST:
 - Complex backgrounds
 - Varying lighting, viewpoints, clutter

These datasets tested whether models could **generalize** to realistic, high-variation images.




ImageNet: The Deep Learning Breakthrough

ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

- Annual competition (2010–2017)
- Millions of images, 1000+ categories

2012 Inflection Point

- **AlexNet** (Krizhevsky, Sutskever, Hinton):
 - Deep CNN, trained on GPUs
- Sparked the modern **deep learning boom** in vision and beyond.



Deep Reinforcement Learning Milestones

Deep Q-Networks (DQN) – 2013

- Mnih et al. (DeepMind)
- Combined: Deep neural nets + Q-learning Learned Atari games directly from **pixels** via trial and error.

AlphaGo – 2016

- DeepMind's system defeated **Lee Sedol**, Go world champion.
- Combined: Deep learning, Monte-Carlo tree search, reinforcement learning
- Showed deep RL can master extremely complex strategy spaces.

- MLPs or CNNs used for classification—the architecture was a direct, continuous mapping from input to output.
 - The Structure: Input → Hidden Layer 1 → Hidden Layer 2 → Output.

The Goal:

Every layer's job was simply to distort, rotate, and scale the data space to make the final classes linearly separable.

limits:

good classification but not for sequence-to-sequence tasks (like translating) or unsupervised representation learning


Encoder–Decoder Architectures

Why encoder–decoder?

- **Encoder:**
 - Compresses high-dimensional input (image, audio, text) into a compact **feature representation**.
- **Decoder:**
 - Expands that representation into a new output format:
 - Pixel-wise segmentation map / Translated sentence / Synthesized audio

Example:

- **U-Net (2015)**: Encoder-decoder CNN for medical image segmentation
 - Produces **pixel-level** output from image inputs.



SECTION II

Generative AI & Large Language Models

- Generating **new** content, not just classifying data
- Multi-modal synthesis (text, image, audio, code)
- GANs, Transformers, and the rise of **LLMs**

What Is Generative AI?

Definition

- Subfield of AI focused on **creating new data**:
 - Natural language text
 - Images and video
 - Audio and music
 - Code and other artifacts

High-level flow

- Massive training data → **GenAI model** → new synthetic outputs
- Model learns **patterns**, not exact copies of training data.

Multi-Modal Output Capabilities

Modern generative models can produce:


- **Audio**
 - Text-to-speech, music composition, voice cloning
- **Video**
 - Short clips from text prompts, trailer generation, restoration
- **Images**
 - Concept art, style transfer, in-painting, design assets
- **Text**
 - Creative writing, summarization, translation, code generation

Same core idea: map a **prompt** → a **new artifact**.

Generative Adversarial Networks (GANs)

Two-network competitive setup:


- **Generator:** Tries to create fake samples that look real.
- **Discriminator:** Tries to distinguish real vs fake samples.



Generative Adversarial Networks (GANs)

Training loop:


1. Generator produces synthetic data.
2. Discriminator scores it.
3. Both networks update based on each other's performance.



Language Models & Embeddings

Language Models (LMs)

- Neural networks trained on large text corpora.
- Goal: model the **probability** of word sequences.




Embeddings

- Map words/tokens → dense numeric vectors.
- Words with similar meaning appear **close** in embedding space:
 - "Grandparent" near "elderly", "adult"
 - "Infant" near "baby", "child"

Embeddings turn discrete language into **continuous geometry**.

Adult, Men, Women,
grandparent
Child, girl, boy,
infant




Context Problems in Static Embeddings

Early embeddings gave **one fixed vector per word**:

- Word “mean” had a single representation, regardless of context:
 - “What do you **mean**?” (verb)
 - “You are being **mean**.” (adjective)
 - “The **mean** absolute error.” (math term)


Without Attention



The Transformer Revolution

Transformers (Vaswani et al., 2017 – “Attention Is All You Need”):

- Replaced recurrence with **attention mechanisms**.
- Key ideas:
 - **Self-attention**: each token can look at **all other tokens**.
 - Learn which parts of the input are most relevant for each position.



Self-Attention: Disambiguating “Bank”

Attention weights connect words to their **context clues**.

Example A (financial):

- “I put my **money** in the **bank**.”
- Attention links “bank” strongly to “money” → financial institution.

Example B (river):

- “The **river** overflowed its **bank**.”
- Attention links “bank” to “river” → geographic feature.

The same word gets **different internal representations** depending on context.

Large Language Models (LLMs)

Characteristics:

- **Billions+** of parameters
- Trained on **hundreds of GBs** of text
 - Equivalent to tens of thousands of full book series

Emergence:

- Above certain scale, models show **unexpected abilities**:
 - Multi-step reasoning
 - Few-shot learning from just a handful of examples
 - Following complex natural language instructions

GPT and Next-Token Prediction

GPT = Generative Pre-trained Transformer

Core mechanism:

- Given previous tokens, predict the **next token**.
- Repeat many times → full paragraph, answer, or dialogue.

Training phases:

1. Pre-training

- Learn general language statistics from massive unlabeled text.

2. Fine-tuning / alignment

- Specialize on tasks, instructions, or safety goals.

From Base Model to Assistant

1. Unsupervised pre-training

- Learn generic language patterns.

2. Supervised fine-tuning

- Train on curated Q&A and instructions.

3. Reward modeling

- Learn what humans prefer in responses.

4. RLHF (Reinforcement Learning from Human Feedback)

- Optimize the model to be:
 - Helpfulm / Honest / Safer

| This turns a raw LLM into a **useful assistant**.

Timeline of LLM Capabilities

- **2018–2019**
 - Early GPT models show promise of Transformers for text.
 - Assistive robots like CIMON support astronauts on the ISS.
- **2020**
 - GPT-3 (175B parameters) demonstrates strong few-shot skills.
 - AlphaFold solves key protein-folding benchmarks.
- **2023+**
 - GPT-4 and multimodal models (text + image + more) become widely available.
 - Millions of users access powerful AI tools daily.

Text-to-Image Generation

Modern diffusion and transformer-based image models:

- Turn detailed **text prompts** into images.
- Control:
 - Style (photorealistic, sketch, anime, etc.)
 - Composition and layout
 - Lighting and atmosphere

Example:

“A cinematic night-time Tokyo street scene, neon lights, rainy reflections”
→ model generates multiple matching images.

Generative Video Capabilities

Recent video models can:

- Generate **short clips** directly from text prompts.
- Maintain **temporal coherence**:
 - Objects persist across frames
 - Motion appears natural
- Approximate simple **physics**:
 - Lighting changes
 - Camera pans and zooms
 - Basic interactions

Potential uses: pre-visualization, VFX, education, rapid prototyping.

Changing Academic Workflows (Expert projections):

- By mid-2020s, advanced AI assistants can:
 - Help explore conjectures
 - Sketch proofs and counterexamples
 - Generate code and experiments

Implications:

- Human researchers focus more on:
 - Problem formulation
 - Conceptual insight
 - Checking and interpreting AI-generated work

AI shifts from **tool** to **collaborator**.

SECTION III

System Safety, Exploits & Vulnerabilities

Hallucinations in LLMs

- LLM outputs that are:
 - Fluent and confident
 - But **factually wrong or invented**

Why it happens:

- Models optimize for **likely text**, not for truth.
- They do not have built-in verified knowledge bases.

Real-world example:

- In *Mata v. Avianca*, a lawyer used ChatGPT to find cases.
- The model fabricated legal citations → court sanctions.

Logic & Reasoning Failure Modes

Pure language models can:

- Follow common text patterns
- But sometimes fail at **precise logical reasoning**

Example: "A juggler juggles 16 balls. Half are golf balls, and half of the golf balls are blue.
How many blue golf balls?"

- Incorrect shortcut: "Half of 16 $\rightarrow$ 8" (ignores the second "half")
- Correct reasoning: 16 balls $\rightarrow$ 8 golf balls $\rightarrow$ 4 blue golf balls.

Modern reasoning-tuned models do better, but the limitation is fundamental:

models track **token probabilities**, not explicit symbolic logic.

Mitigating Hallucinations in Practice

Defensive strategies:

- **Contextual constraints**
 - System prompts with explicit rules and boundaries.
- **Conversation framing**
 - Keep the model within a known domain or document scope.
- **Sampling controls**
 - Lower “temperature” or adjust decoding to reduce wild outputs.
- **Post-processing**
 - Sanity checks, format validation, or secondary verification tools.

| Goal: encourage **grounded**, not speculative, answers.

Static Knowledge vs Dynamic Reality

Limitations of static training: After training, an LLM's internal knowledge is **frozen** at a **cutoff date**.

- It cannot:
 - Natively know post-cutoff events
 - Provide sources or citations for its statements.

Trust issues:

- Without external grounding, users cannot easily:
 - Verify where an answer came from
 - Check if it is up to date

| This motivates hybrid architectures like **RAG**.

Retrieval-Augmented Generation (RAG)

How RAG works:

1. **User query:** e.g., "Summarize the latest policy from document X."
2. **Retrieval:** Search external databases / documents for relevant passages.
3. **Augmented prompt:** Combine user question + retrieved context.
4. **Generation:** LLM answers *using that specific context*.
5. **Citations:** System can point back to retrieved sources.

Benefit:

- Answers are **grounded** in actual documents, reducing hallucinations and updating knowledge without retraining.

Adversarial Security Exploits (Intro)

LLM systems also face **security threats**:

- **Prompt injection**: User tries to override system instructions.
- **Jailbreak attempts**: Crafted prompts to bypass safety filters.
- **Data exfiltration**: Extracting confidential information from integrated systems.

Mitigation requires:

- Careful system prompts and guardrails
- Input/output filtering
- Separation of sensitive tools and data

| *(Deeper security patterns can be expanded in a dedicated session.)*

Security Threats: Prompt Injections & Jailbreaking

As LLMs are wired into tools and workflows, new security risks appear:

- **Prompt injection attacks:** Malicious text tries to override system instructions and safety rules.
- **Indirect injection:** Hidden instructions embedded in web pages / documents that the LLM reads.
- **Common attack patterns**
 - (1) Jailbreak prompts, (2) Context manipulation ("ignore previous instructions..."), (3) Obfuscated code or data to exfiltrate secrets, (4) Attempts to extract credentials or internal configuration

Defensive design is as important as model accuracy.

SECTION IV

Privacy, Trust & Governance

The Myth of Simple Anonymization

Removing obvious identifiers (name, ID) is **not enough**:

- **Linkage attacks**
 - Cross-link “anonymous” data with public sources to re-identify people (e.g., Netflix Prize dataset deanonymization).
- **A-priori knowledge**
 - Attacker uses a few known facts about someone to locate their record.
- **Composition attacks**
 - Combine multiple releases over time to triangulate identities.

Conclusion: privacy must be **designed**, not assumed.

Federated Learning: Decentralized Training

Goal: train a **shared model** without centralizing sensitive data.

- **Local training**
 - Each site (e.g. hospital, bank branch) keeps data on-prem.
 - Trains a model update on its own data.
- **Central aggregation**
 - Only **model updates / weights** are sent to a server.
 - Server aggregates into a global model.

Federated learning improves privacy but does **not** solve it alone.

Differential Privacy: Math-Backed Protection

Differential Privacy (DP) gives **formal guarantees**:

- **Core idea**
 - Whether or not any single person's data is in the training set should barely change the model's outputs.
- **Mechanism**
 - Add carefully calibrated **noise** to: Data | Gradients | Query answers

Example:

- **Randomized response**
 - Individuals sometimes flip answers according to a known probability.
 - Protects individuals while preserving accurate **aggregate statistics**.

Model Inversion & Reconstruction Risks

Without privacy controls, models can **memorize** data:

- **Reconstruction attacks**
 - Adversary probes a trained model to reconstruct training examples:
 - Faces, text snippets, health records, etc.
- **Membership inference**
 - Attacker infers whether a specific person's data was in the training set.

Mitigation:

- Train with **differential privacy** and regularization.
- Limit access to model internals and predictions.

Explainable AI (XAI)

Why we need explanations: (1) High-stakes decisions (credit, hiring, medicine, justice) (2) require **justification**, not just accuracy.

Approaches:

- **Global explanations:** Which features matter overall?
- **Local explanations:** Why this particular prediction?

Example method:

- **Feature attribution** (e.g. RISE-style masking)
 - Randomly mask parts of an input (pixels, words)
 - See how prediction changes
 - Highlight the regions that contribute most to the decision.

The Future of Autonomous Agents

Beyond static chat:

- **Traditional LLMs**
 - Reactive: answer questions, translate, summarize.
- **Agentic LLMs**
 - Can:
 - Call tools and APIs
 - Browse the web
 - Write and run code
 - Plan and execute multi-step tasks

The Future of Autonomous Agents

Research example:

- **Stanford small-town simulation**
 - Dozens of LLM-powered agents lived in a sandbox world.
 - They independently:
 - Planned a party
 - Invited others
 - Coordinated schedules
 - Held the event
 - Demonstrated emergent social behavior.

Prompt Engineering: Interacting with LLMs

Prompt engineering = designing prompts to steer model behavior without changing model weights.

High-level flow:

- **Target objective + system context**
↓
- **Optimized prompt** (instructions, examples, format)
↓
- **High-precision output** (more reliable and on-task)

Good prompts act as a **user-level programming language** for LLMs.

Core Prompting Strategies

Common patterns:

- **Basic task prompts**
 - Summarization, classification, translation, role-playing.
- **Zero-shot prompting**
 - Ask directly, no examples.
- **Few-shot prompting**
 - Include a handful of input–output examples in the prompt.
- **Chain-of-Thought (CoT)**
 - Tell the model: “Think step by step.”
 - Encourages explicit reasoning → often **better accuracy** on logic problems.

Conclusion & Core Takeaways

- **Paradigm shift**
 - From hand-coded rules → data-driven optimization and learning.
- **Infrastructure realities**
 - AI is bounded by data volume, network bandwidth, and compute hardware.
- **Safety & governance**
 - Hallucinations, privacy leaks, and prompt injections require careful system design and monitoring.
- **Path forward**
 - Combine **scale and capability** with:
 - Privacy and security
 - Explainability and governance
 - Thoughtful human oversight

